



## DATA SCIENCE 102: INTRODUCTION TO MACHINE LEARNING

# AGENDA

---



- What is Machine Learning?
- Types of Machine Learning
  - Applications of Machine Learning
- General Machine Learning Steps using Sci-Kit Learn
  - Fundamentals
    - Standard steps
    - Reading the library
  - Examples
    - Regression

# WHAT IS MACHINE LEARNING


---



# WHAT IS MACHINE LEARNING



What Society Thinks I do



What My Friends Think I Do



What other Computer Engineers think I do

```
In [1]:  
import keras  
Using TensorFlow backend.
```

What I actually do



# WHAT IS MACHINE LEARNING

---



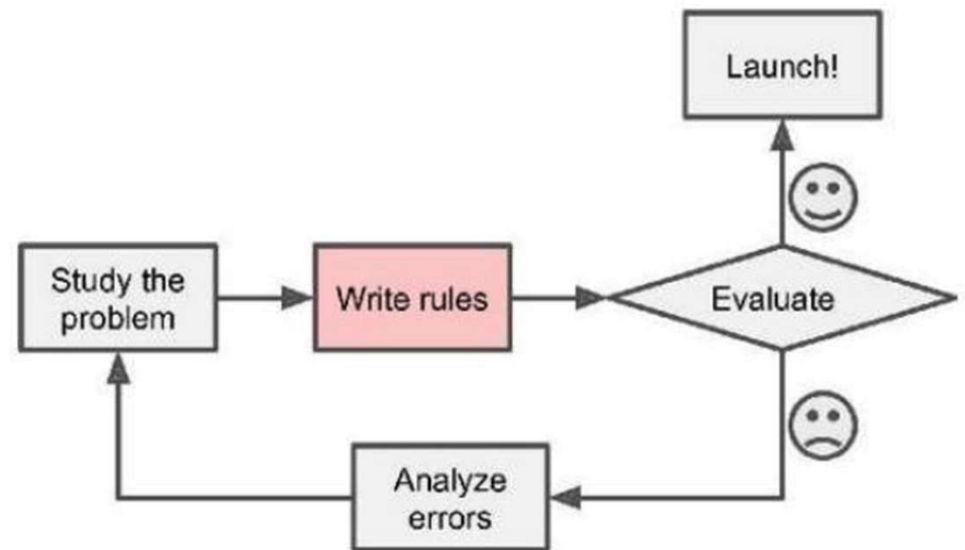
A study of algorithms and statistical models  
that computers use to perform a specific task effectively  
without being explicitly programmed

# EXAMPLE - SPAM EMAIL DETECTION



## Without Machine Learning

- You will write explicit rules to filter out emails containing words like:
  - "4U", "Loan", "Prince", "Money Transfer", "Free"
- Classify any email that contain these text as spam email

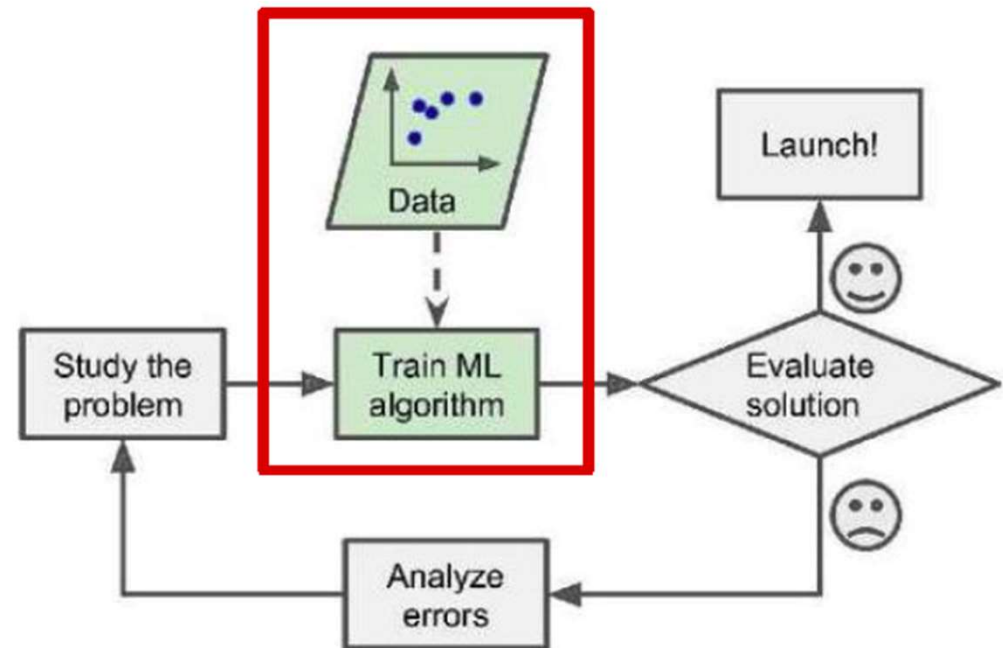


# EXAMPLE - SPAM EMAIL DETECTION



## With Machine Learning

- You train a model on some train data (spam & non-spam emails) with an algorithm that **infers** characteristics and patterns of spam emails
- The trained model can then apply learnt rules to detect signals of spam in an incoming email





# TYPES OF MACHINE LEARNING

---

- Types of Machine Learning
- Supervised Applications
- Unsupervised Applications



HACKWAGON  
• ACADEMY •





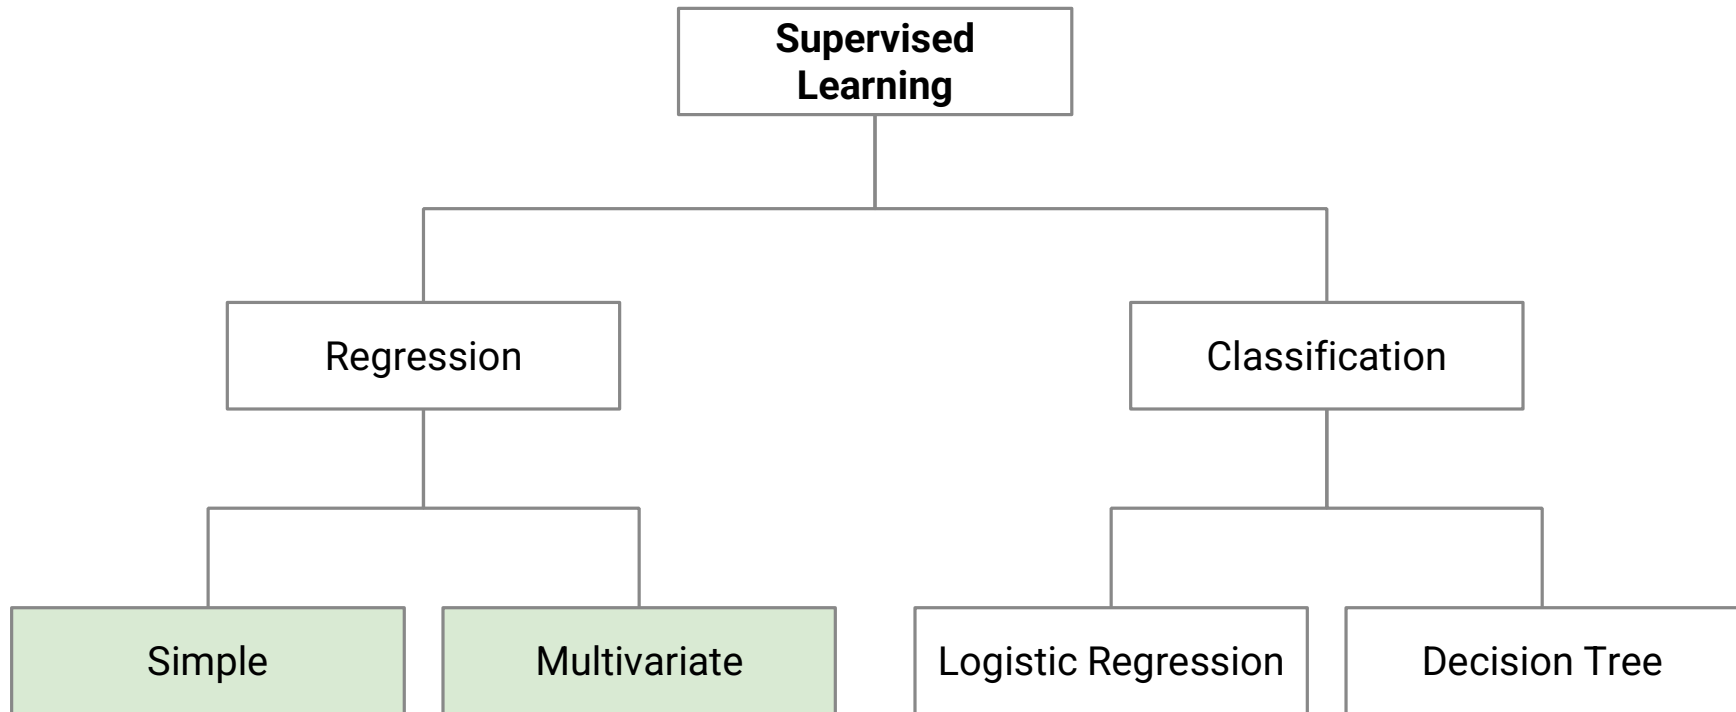
# TYPES OF MACHINE LEARNING

---

- Supervised Learning
  - Based on **labelled data**, makes **predictions** on a test set
- Unsupervised Learning
  - Based on **unlabelled data**, algorithm **discovers patterns** within the dataset

# SUPERVISED LEARNING


---



# SUPERVISED APPLICATIONS



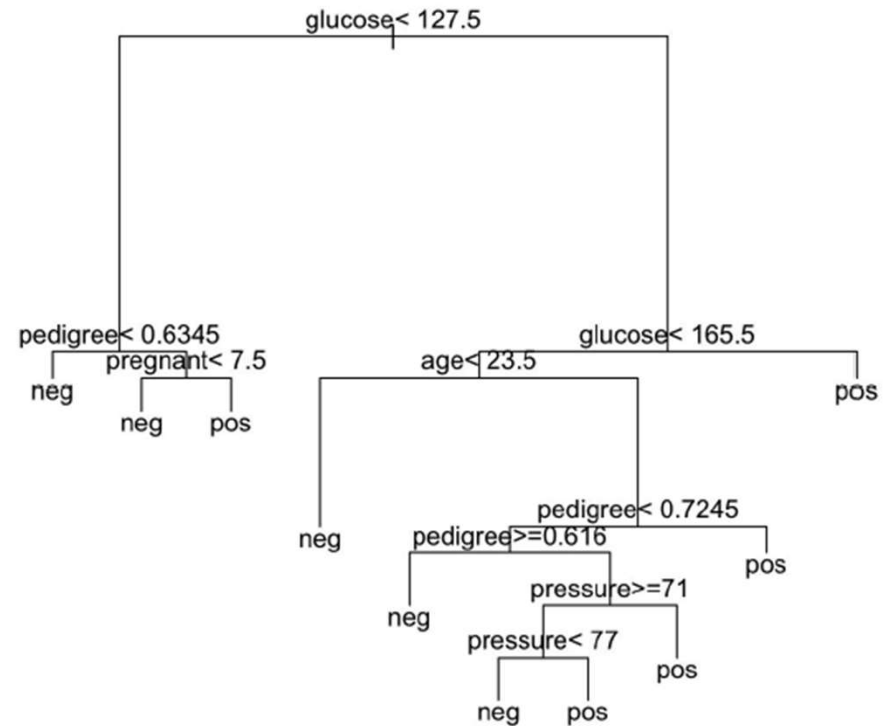
- Regression
  - Predict Housing Price (Target Variable) from Housing Floor Area (Input Feature)



# SUPERVISED APPLICATIONS




- Classification
  - Predict if a patient has diabetes (Yes/No)



# UNSUPERVISED APPLICATIONS



- Clustering
  - Given a set of unlabelled data points, identify cluster each point belongs to



## GENERAL MACHINE LEARNING STEPS USING SKLEARN

---

- Fundamentals
  - Standard Steps
  - Reading the library
- Examples
  - Regression





## FUNDAMENTALS - STANDARD STEPS

---

When using any machine learning models with scikit learn, the following steps are usually applied in order:

**Step 1:** Choose a class of machine learning model from the library

**Step 2:** Choose the model's hyperparameters by instantiating with desired values (tuning)

**Step 3:** Arrange data into features and target

**Step 4:** Fit model to your data by using the `fit()` method of the model

**Step 5:** Apply the model to new data:

- For supervised learning, using the `predict()` method
- For unsupervised learning, using the `predict()` or `transform()` method





# FUNDAMENTALS - READING THE LIBRARY

- Determine what model is used, then delve into details of that model in the library, and the other required parts of the library
- Here are some of the common classes you would use:
  - `sklearn.<a name of a model>` (i.e `sklearn.linear_model`)
    - Access to all functionalities of that model
  - `sklearn.metrics`
    - Functionalities of assessing the performance of your models
  - `sklearn.model_selection`
    - Performs cross validation and tuning of parameters
  - `sklearn.feature_selection`
    - Reduce dimension of dataset to boost performance



# EXAMPLE - REGRESSION

```
from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score
```

Step 3

```
# Load the diabetes dataset
diabetes = datasets.load_diabetes()

# Use only one feature
diabetes_X = diabetes.data[:, np.newaxis, 2]

# Split the data into training/testing sets
diabetes_X_train = diabetes_X[:-20]
diabetes_X_test = diabetes_X[-20:]

# Split the targets into training/testing sets
diabetes_y_train = diabetes.target[:-20]
diabetes_y_test = diabetes.target[-20:]
```

Step 1 & 2

```
# Create linear regression object
regr = linear_model.LinearRegression()
```

```
# Train the model using the training sets
regr.fit(diabetes_X_train, diabetes_y_train)
```

Step 4

```
# Make predictions using the testing set
diabetes_y_pred = regr.predict(diabetes_X_test)
```

Step 5



## DATA SCIENCE 102: CLASSIFICATION

# AGENDA

---



- What is Classification
- Logistic Regression
- Confusion Matrix



## LEARNING OBJECTIVES

---

- Build an intuition for Logistic Regression
- Able to interpret results and performance from Logistic Regression
- Apply Confusion Matrix to evaluate classification model

# LOGISTIC REGRESSION

---

- Intuition
- Confusion Matrix





# Logistic Regression

---

- Logistic Regression is one of the basic and popular algorithms to solve a binary classification problems
- For each input, logistic regression outputs a probability that this input belongs to the 2 classes
  - Set a probability threshold boundary and that determines which class the input belongs to
- Binary classification problems (2 classes):
  - Emails (Spam / Not Spam)
  - Transactions (Fraudulent / Not Fraudulent)
  - Loan Default (Yes / No)





# Why Not Use Linear Regression for Classification?


---

- Linear Regression is used to solve regression problems with continuous values
- Logistic Regression is used to solve classification problems with discrete categories
  - Binary classification (Classes 0 and 1)




# Why Not Use Linear Regression for Classification?

- Let's say we want to predict the malignancy of a tumour based on its size
- We then plot a simple linear regression line and set the threshold as 0.5
  - Negative class – Tumour size on the left size
  - Positive class – Tumour size on the right size





Another example: Negative probability





# Why Not Use Linear Regression for Classification?


- Solution – Transform linear regression to a logistic regression curve
- Logistic regression is a Sigmoid function
  - Sigmoid function takes in any real value
  - Output probability between 0 and 1





# Sigmoid Function (Logistic Function)



- Take linear regression and put in into the Sigmoid function
- Sigmoid function outputs probability between 0 and 1





# Sigmoid Function (Logistic Function)

- Sigmoid function outputs probability between 0 and 1
- Default probability threshold is set as 0.5 typically
  - Class 0 – Below 0.5
  - Class 1 – Above 0.5





## Model Evaluation – Confusion Matrix

- To evaluate logistic regression model, confusion matrix is typically used for model evaluation
- Imagine now a doctor wants to use a logistic regression model to predict if patients have (or don't have) disease

n=165	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	



- True Positives (TP)
- True Negative (TN)
- False Positives (FP)
- False Negatives (FN)

# Metrics for Model Evaluation is Classification



- Accuracy

➤  $\frac{TP+}{Total}$

- Precision

➤  $\frac{TP}{TP+}$

- Recall

➤  $\frac{TP}{TP+FN}$

n=165	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	

Precision:  $100/110 \rightarrow 90.9\%$

Recall:  $100/105 \rightarrow 95.2\%$

## Example - Disease Detection

- We want to focus on **Recall** as the main metric
- Reduce FN and increases Recall
- Because if an infected patient goes through the test undetected, the risk is very high if the disease is contagious





## DATA SCIENCE 102: REGRESSION

# AGENDA

---



- What Is Regression
  - Applications of Regression
- Simple Linear Regression
  - Intuition
  - Interpretation
  - Ordinary Least Squares
  - Measure of Fit
  - Model Validation - Train Test Split
  - Performance Measures
  - Scikit Learn Example and Practice
- Multivariate Linear Regression
  - Interpretation
  - Feature Selection
  - Scikit Learn Example and Practice



## LEARNING OBJECTIVES

---

- Build an intuition for Linear Regression
- Able to interpret results and performance from Linear Regression
- Discern difference between Single and Multivariate Linear Regression
- Importance of Feature Selection in Multivariate Linear Regression

# WHAT IS REGRESSION

---

- Applications of regression





# WHAT IS REGRESSION

---

- Regression Analysis is a model to develop an equation that shows how variables are related
- In regression terminology, the variable being predicted is the **dependent variable** and the variables being used to predict the dependent variable are called **independent variables** (predictor variables)
- Regression can be used in the following cases:
  - Predicting continuous labels
  - Classification (logistic regression)

# SIMPLE LINEAR REGRESSION

---

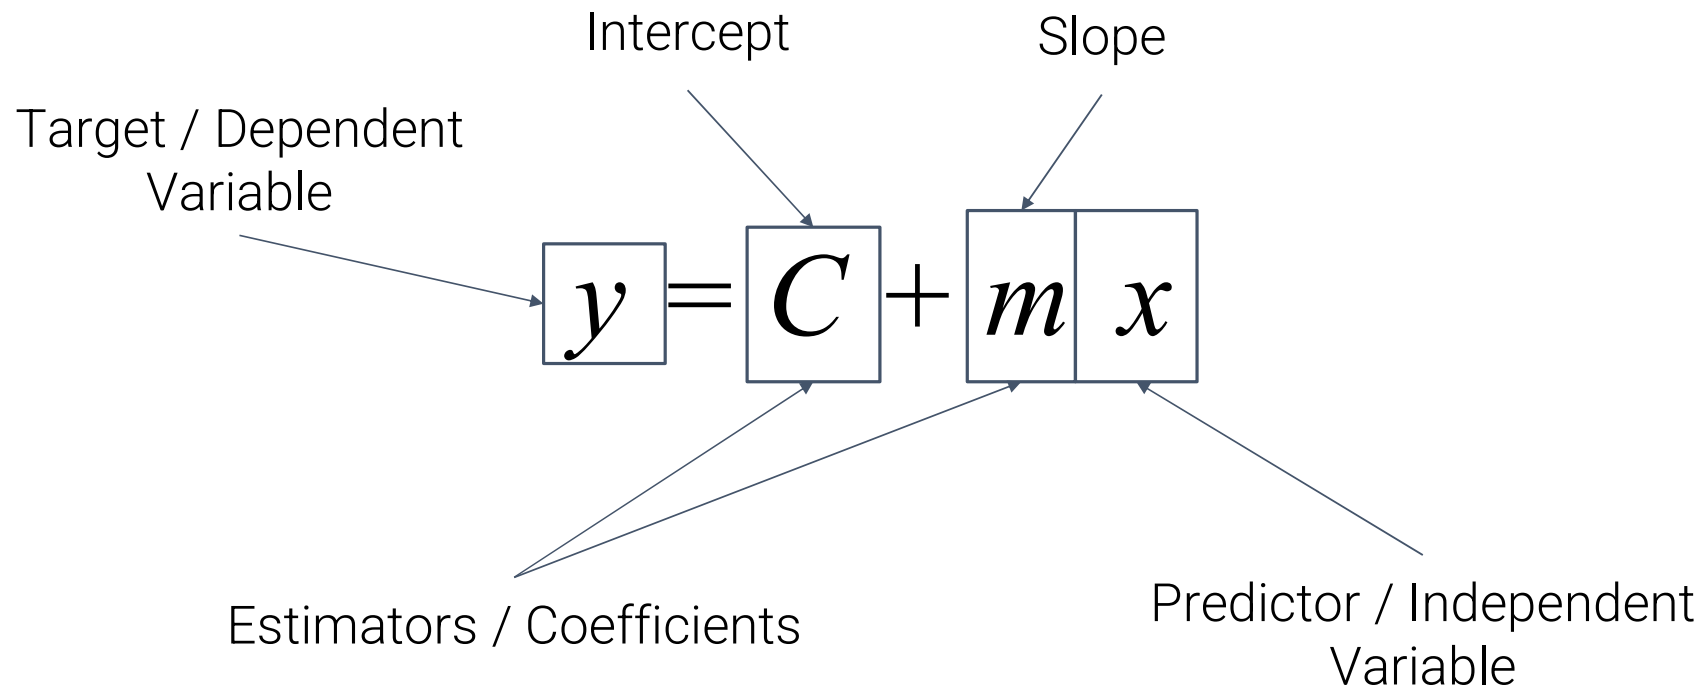
- Intuition
- Single vs Double variable
- Interpretation
- Ordinary Least Squares
- R Squared and Adjusted R Squared
- Performance Measures
- Train-Test-Split





# INTUITION - THE REGRESSION EQUATION

Remember the  $y = mx + c$  you learnt in secondary school? That is essentially regression!

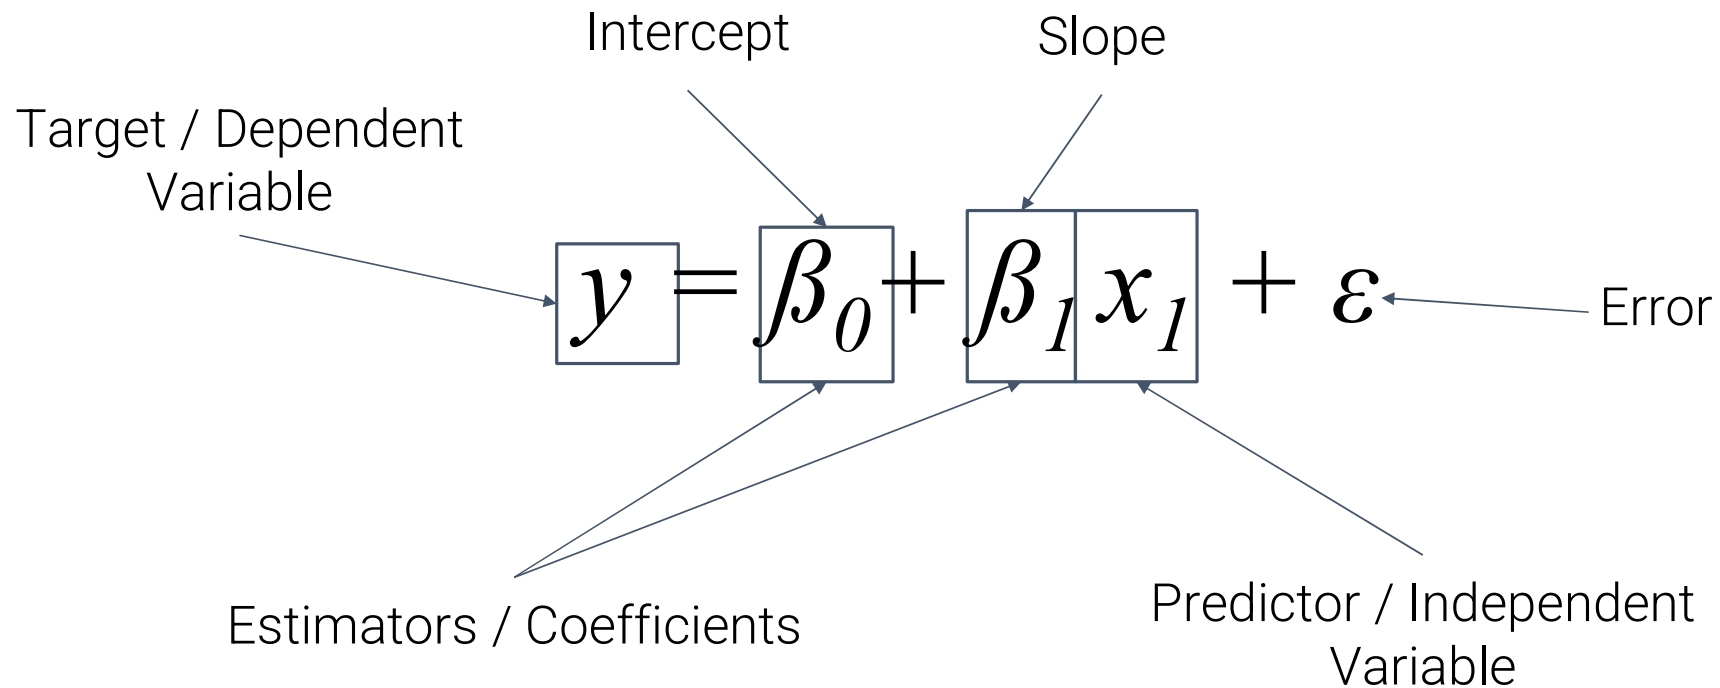






## INTUITION - THE REGRESSION EQUATION

This is the adult, and machine-learning way of representing regression:



# STUDY OF SINGLE VARIABLE - NO VALUE ON ITS OWN




- Predict future tips given the following dataset:
  - To “predict” future tips, use average

$$\text{Tip} = 0 \text{ (No Predictor)} + 15$$

$$\text{Tip} = 15$$

Tip #	Amount (\$)
Tip 1	10
Tip 2	25
Tip 3	12
Tip 4	8
Tip 5	15
Tip 6	20




# STUDY OF TWO VARIABLES - SIMPLE LINEAR REGRESSION



- Predict future tips given the following dataset
  - To “predict” future tips using linear regression

Tip #	Amount (\$)	Meals (\$)
Tip 1	10	80
Tip 2	25	150
Tip 3	12	75
Tip 4	8	60
Tip 5	15	100
Tip 6	20	150

$$\text{Tip} = -1.27 + 0.158 (\text{Meal})$$





## INTERPRETATION

---

- After fitting the model, you can extract the coefficients (estimators) of the model by using `<variablenameofmodel>.coeff_`. To get the intercept, use, `<variablenameofmodel>.intercept_`
- Let's say given the following regression equation:

$$Tip = - 1.27 + 0.158 (Meal)$$

- We can interpret it as, given **1 dollar increase** in the cost of a meal, there will be an **increase** of tips by **0.158 dollars**



## INTERPRETATION - CODED EXAMPLE

- Let's say given the following regression equation:

$$\text{Tip} = -1.27 + 0.158 (\text{Meal})$$

```
1 print(regr.coef_)
2 print(regr.intercept_)
3
4 # Tips = -1.27 + 0.158 (Meal)
5 # With every $1 increase in meal, tips would increase by $0.15
```

```
[[0.15881384]]
[-1.27841845]
```

- So what happens behind `.fit()` == OLS



## HOW LINEAR REGRESSION EQN IS CONSTRUCTED - OLS

- Now we understand Linear regression is about creating an equation that shows how variables are related. In code form your “.fit()” will create this equation.
- This equation is useful in helping us (1) interpret relationship, and (2) predict output based on unseen input data

Means we can sub in any X-Value (meal) to predict tips

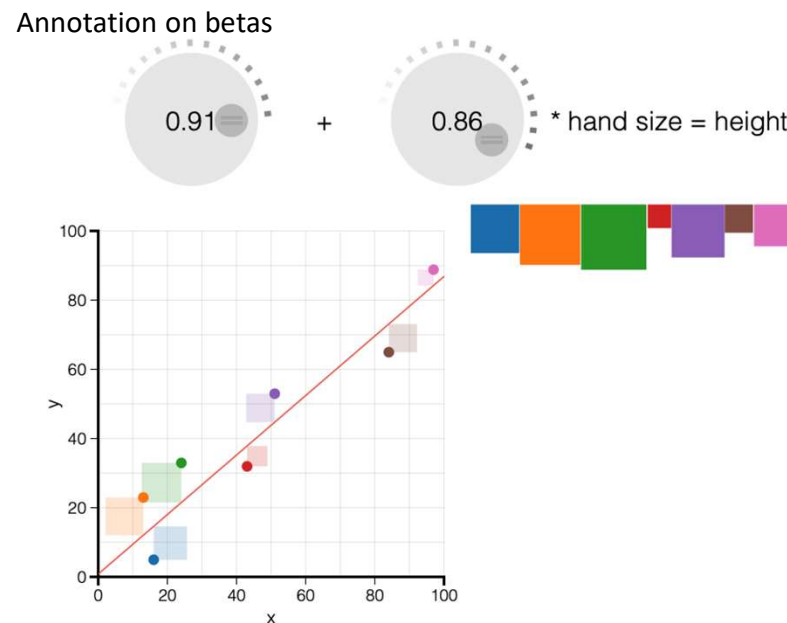
$$Tip = -1.27 + 0.158 (Meal)$$

- Then the question is, how do we derive the y-intercept ( $\beta_0$ ) and gradient ( $\beta_1$ )? Answer is: Ordinary Least Squares (OLS).

# ORDINARY LEAST SQUARES - VISUALISATION OF PROCESS



- Ordinary Least Squares (Least Squares Method) is aimed at **minimizing** the sum of squared residuals (SSR), i.e keep the errors ( $\epsilon$ ) as low as possible



Click here to view a visualisation of OLS in action: <http://setosa.io/ev/ordinary-least-squares-regression/>



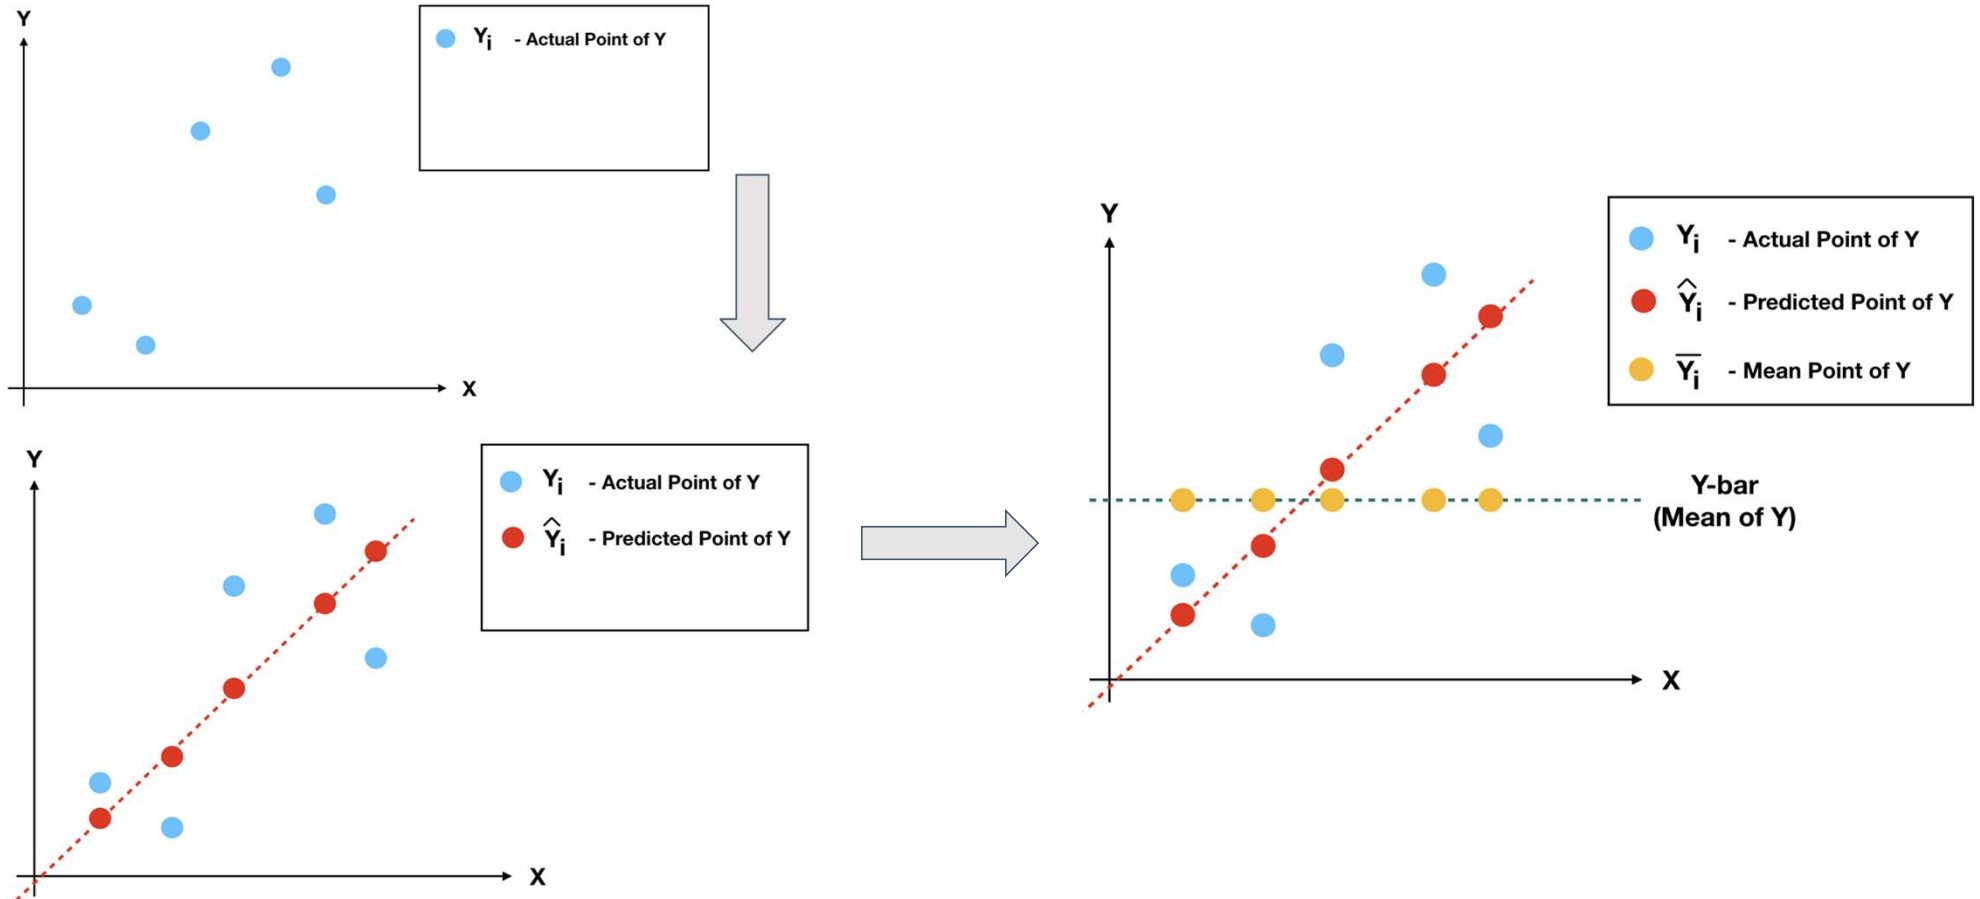
## ORDINARY LEAST SQUARES - KEY CONCEPT

---


- Intuitively, OLS is fitting a line through the sample points such that the sum of squared errors (SSR) is as small as possible, hence the term least squares
- The OLS regression line always goes through the mean of the sample ( $\bar{y}$ )
- The residual,  $e$ , is an estimate of the error term,  $\varepsilon$ , and is the difference between the fitted line (sample regression function) and the sample point
- The sum of the OLS residuals is zero



# ORDINARY LEAST SQUARES - BUILDING BLOCKS EXPLAINED



# ORDINARY LEAST SQUARES - BUILDING BLOCKS EXPLAINED



## MEASURING FIT BETWEEN VALUES AND PREDICTED LINE - $R^2$



- A measure of how well the sample values fit the predicted regression line is **goodness of fit**
- $R$  (also known as Pearson's  $R$ ) is the sample correlation coefficient; ranges from -1 to 1;  $>0$  = positive correlation,  $<0$  = negative correlation
- $R^2$  (R-squared) is known as the coefficient of determination; range from 0 to 1

$$\frac{SSE}{SSR + SSE} = \frac{SSE}{SST} = R^2$$

*An Introduction to Statistical Learning (with applications in R) by Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani (2013)*

# CALCULATION OF $R^2$



- Each observation ( $y_i$ ) on the linear regression is made up of an explained part ( $\hat{y}_i$ ) and an unexplained part ( $e_i$ ):


$$y_i = \hat{y}_i + e_i$$

Total sum of squares (SST) =  $\sum (y_i - \bar{y})^2$

Explained sum of squares (SSE) =  $\sum (\hat{y}_i - \bar{y})^2$

Residual sum of squares (SSR) =  $\sum (y_i - \hat{y}_i)^2$

$$\text{SST} = \text{SSE} + \text{SSR}$$



*An Introduction to Statistical Learning (with applications in R)* by Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani (2013)

## AN ALTERNATIVE OF MEASURING FIT - ADJUSTED $R^2$



- Adjusted  $R^2$  modifies the original  $R^2$  by incorporating the sample size and the number of explanatory variables in the model
- Can be found in `sklearn.metrics.r2_score` or `LinearRegression.score()`

$$R_{adj}^2 = 1 - \left[ \frac{(1 - R^2)(n - 1)}{n - k - 1} \right]$$



## R-SQUARED VS. ADJUSTED R-SQUARED

---

- Both  $R^2$  and the adjusted  $R^2$  (Adj- $R^2$ ) gives an idea of how many data points fall within the line of the regression equation
- The main difference is:
  - $R^2$  tells you every single variable explains the *variation* in the depend variable ( $y$ )
  - Adj- $R^2$  tells you the **percentage of variation** explained only by the independent variables that actually affect the dependent variable

# R-SQUARED VS. ADJUSTED R-SQUARED



- Formula for Adj- $R^2$  :

$$\text{Adjusted } R^2 (R^2_{\text{adj}}) = 1 - \left[ \frac{(1-R^2)(n-1)}{n-k-1} \right]$$

- $n$  is the number of points in your data sample
- $k$  is the number of independent variables (predictors) excluding constants



## R-SQUARED VS. ADJUSTED R-SQUARED

---

- The adjusted  $R^2$  (Adj- $R^2$ ) penalizes for adding more independent variables that do not fit the model.
- As  $R^2$  increases with every predictor added to a model, it can appear to be a **better fit** with more terms added to the model, but this can be misleading
- Adding too many variables and polynomials may run into the trouble of **overfitting**, thus a misleading high  $R^2$  value can lead to misleading projections
- In short, adjusted  $R^2$  is preferred over  $R^2$  in multivariate linear regression.



# MEASURING FIT VS MEASURING PERFORMANCE

---



- Measuring Fit - Quantified via  $R^2$  / Adjusted  $R^2$ , is a measure of how well the seen / known values fit the predicted regression line
- Measuring Performance - is a measure of how well the predicted regression line can predict unseen/unknown values (Quantified RMSE or MAE)



## PERFORMANCE MEASURES - LINEAR REGRESSION

- For linear regression there are several measures to assess the performance of the model:
  - **MAE** (`sklearn.metrics.mean_absolute_error`)
    - Mean Absolute Error
    - Whether predictions are, on average, over/under predicting the outcome
  - **RMSE** - (`sklearn.metrics.mean_squared_error` then use `** 0.5`)
    - Root Mean Squared Error
    - Differences between predicted vs observed/actual values
    - Similar to standard error
    - Lower is better
    - 0 means perfect fit to data (*not the best, could be overfitting*)
- Can be found in the `sklearn.metrics` part of `sklearn`

*An Introduction to Statistical Learning (with applications in R)* by Gareth James, Daniela Witten, Trevor Hastie and Robert Tibshirani (2013)

# PERFORMANCE MEASURES - LINEAR REGRESSION

---




- How are RMSE and MAE derived?
- Answer: Via Model-Validation (Train-Test-Split)



## MODEL VALIDATION - TRAIN TEST SPLIT

- After selecting and fitting a machine learning model, like Linear Regression for example, we need to ensure that we validate our model **by comparing some of the training data and comparing the prediction against its known value**
- One of the ways to validate the model would be to use *Holdout Sets*, basically splitting the given datasets for training and testing
- We can use Scikit learn's `train_test_split` to split our datasets
- By default, it is an 80-20 split (80% training, 20% testing)




# MODEL VALIDATION - TRAIN TEST SPLIT - CODED EXAMPLE



X1... Xn	Y
X Training	Y Training
X Testing	Y Testing

```
3  
4 housing_x = x_all[['sqft_living', 'floors']]  
5  
6 X_train, X_test, y_train, y_test = train_test_split(housing_x, y_all, random_state=42)
```



## SCIKIT LEARN EXAMPLE AND PRACTICE\*

---



- Try out the practice in your in-class notebook 5
- Remember the 5 common steps for using SKLearn

# MULTIVARIATE LINEAR REGRESSION

---

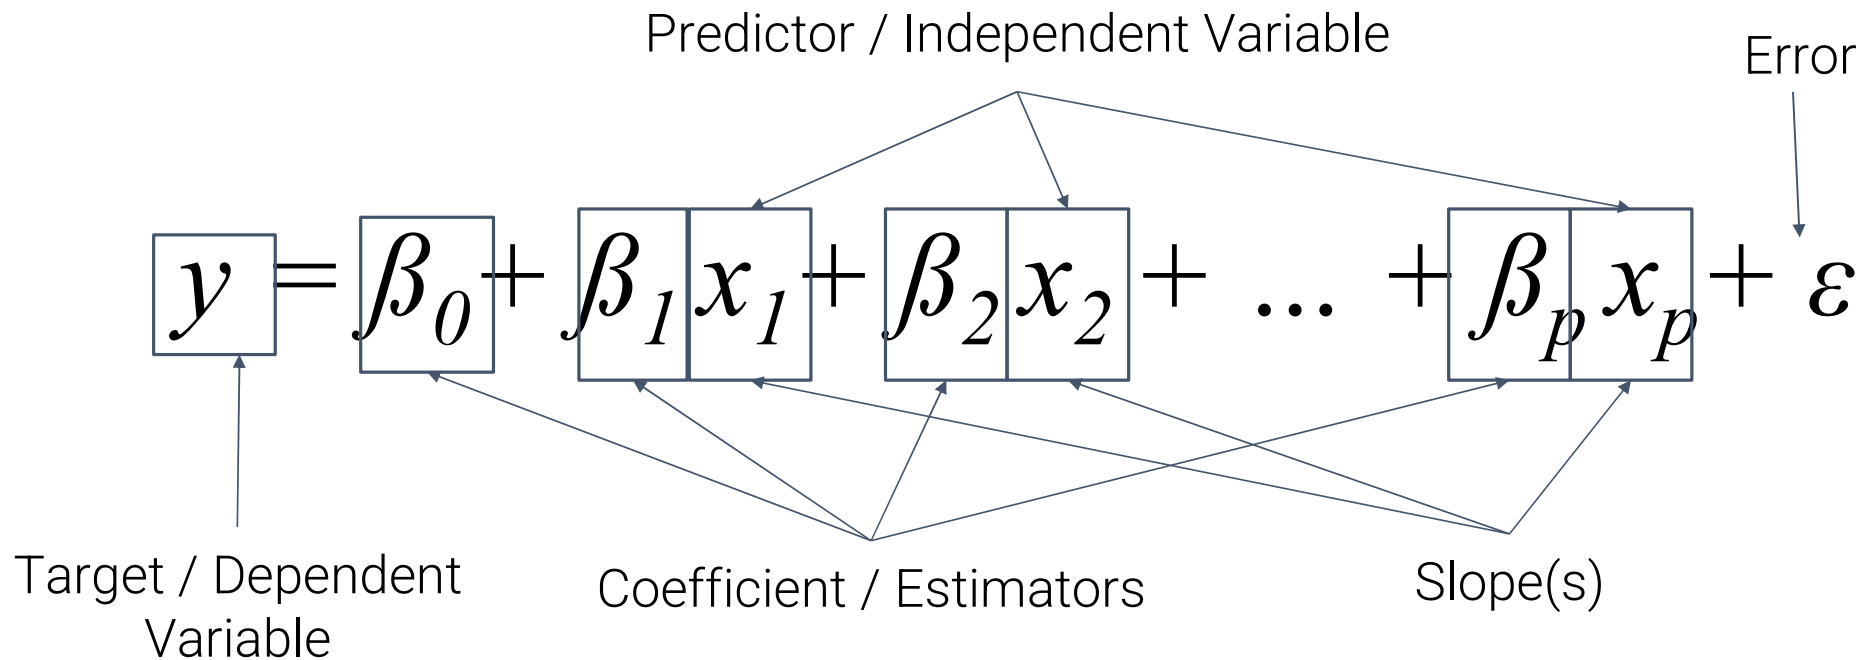
- Intuition and Interpretation
- R-Squared vs. Adjusted R-Squared
- Feature Selection



# INTUITION



- Multivariate linear regression incorporates **more than one predictors** into the equation





## INTERPRETATION

---



- Let's say given the following regression equation:

$$\text{HousingPrice } (\$) = - 48237.317 + 274 \text{ SQFT} + 12566.866 \text{ Floor} + \varepsilon$$

- We can interpret it as, given 1 SQFT increase, Housing Price increases by \$274; given 1 floor increase, Housing Price increase by \$12566.866



## INTERPRETATION - CODED EXAMPLE

- Let's say given the following regression equation:

$$\text{HousingPrice (\$)} = -48237.317 + 274 \text{ SQFT} + 12566.866 \text{ Floor} + \varepsilon$$

```
1 print(multi_housing_lr.coef_)
2 print(multi_housing_lr.intercept_)
3
4 # Housing Price = -48237.317 + 274 (SQFT_Living) + 12566.866 (Floors)
```

```
[[ 274.0203467 12566.86687756]]
[-48237.31783364]
```

# FEATURE SELECTION

---



- Why don't we just use all the variables in the world and just apply it to our model?
  - Expensive or not feasible
  - Sometimes fewer predictors are better
  - More predictors could lead to possibly more missing data
  - Lesser predictors allow for greater insight into "influence"
  - Unstable regression coefficient due to multicollinearity
- Approach to reducing / selecting predictors:
  - Domain expert eliminate irrelevant predictors
  - Summary statistics - Frequency and correlation plots

## FEATURE SELECTION - MULTICOLLINEARITY



- Multicollinearity occurs when **one predictor variable** in a model can be linearly predicted with other **predictor variables**.

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \varepsilon$$

The diagram illustrates multicollinearity by showing three predictor variables,  $x_1$ ,  $x_2$ , and  $x_p$ , each enclosed in a box within the regression equation. Three double-headed arrows connect these boxes, indicating linear relationships between them. The arrow between  $x_1$  and  $x_2$  is labeled "Correlated" below it. The arrow between  $x_1$  and  $x_p$  is labeled "Correlated" below it. The arrow between  $x_2$  and  $x_p$  is labeled "Correlated" above it.

# FEATURE SELECTION - MULTICOLLINEARITY

---



- Consequences of this issue includes:
  - Loss of precision
  - $R^2$  value takes on a high value despite not being statistically significant
  - Some of the signs of the coefficient might change
- Remedies to this problem includes:
  - Dropping problem variables (*selecting which feature to drop*)
  - Remedy using domain expertise
  - Get more data

## FEATURE SELECTION - MULTICOLLINEARITY INTUITION



- Given the following dataset, there is a clear multicollinearity between the two predictor variables


House (#)	Price (\$)	Price per square <b>foot</b> (\$)	Price per square <b>metre</b> (\$)
1	10000	80	240
2	25000	150	450
3	12000	75	225
4	8000	60	180
5	1500	100	300
6	2000	150	450

$$Price = \beta_0 + \beta_1 SQFT + \beta_2 SQM + \varepsilon$$

# FEATURE SELECTION - DETECTING MULTICOLLINEARITY



- You can detect the correlation between variables using a correlation matrix (Rule of thumb: correlation coefficients  $> 0.8$  signals multicollinearity)



# FEATURE SELECTION - DETECTING MULTICOLLINEARITY

---



- Other ways of detecting multicollinearity include:
  - Variance Inflation Factor
  - Low Variance
  - `sklearn.feature_selection.f_regression`



## SCIKIT LEARN EXAMPLE AND PRACTICE\*

---



- Try out practice 2 in your in-class notebook
- Make sure you are able to identify the features which are correlated

# THANK YOU!

---



HACKWAGON  
• ACADEMY •